

1 Temporal Interview Profiling System (TIPS)

1.1 Automated Interview Analysis System

- Processes pre-recorded video interviews
- Extracts behavioral signals from audio and video
- Evaluates candidate responses using LLMs
- Generates time-evolving performance scores
- Compares answers against job descriptions

2 Abstract

- TIPS is an automated interview analysis system for video interviews
- Multi-stage backend pipeline processes recordings through 6 stages
- Extracts audio signals: transcription, pitch, energy, speech rate
- Analyzes video signals: face detection, head pose, gaze stability
- Uses Qwen2.5-3B LLM for semantic relevance scoring

3 Problem Statement

- **Subjectivity:** Human interviewers have inherent biases and varying standards
- **Limited Analysis Depth:** Manual evaluation misses vocal and temporal patterns
- **Scalability Issues:** Time for evaluations grows linearly with candidate volume
- **Lack of Standardization:** Comparing candidates across interviewers is difficult
- **No Temporal Insights:** Traditional scoring gives only aggregate scores

4 Motivation

- Modern recruitment faces bottlenecks in interview evaluation
- Organizations invest significant resources with inconsistent results
- Automated systems can capture signals humans often miss
- Modular design allows independent optimization of each stage
- Functions as addon to existing interview platforms

5 System Architecture

5.1 Pipeline Flow

1. **Input:** Interview Recording + Job Description
2. **Stage 0:** Timebase Establishment
3. **Stage 1:** Parallel Signal Extraction (Audio/Video)
4. **Stage 2:** Temporal Segmentation (Q&A Pairing)
5. **Stage 3:** Behavioral Metrics Computation
6. **Stage 4-5:** LLM Semantic Scoring + Verdict

5.2 Components

- Interview UI Layer (WebRTC Recording)
- Backend Pipeline (6-Stage Processing)
- Dashboard (Visualization - Planned)

6 Methodology

6.1 Pipeline Overview

- **Stage 0:** Establishes canonical timebase from video metadata
- **Stage 1:** Parallel extraction of audio features, video features, transcription
- **Stage 2:** Identifies speaking segments and pairs Q&A
- **Stage 3:** Computes behavioral metrics (confidence, fluency, eye contact)
- **Stage 4-5:** LLM scoring and final verdict aggregation

6.2 Key Processing Stages

- **Faster-Whisper:** Speech-to-text with word-level timestamps
- **WebRTC VAD:** Voice activity detection for speech segments
- **MediaPipe:** Face detection, landmark tracking, head pose estimation
- **Librosa:** Audio features (pitch, energy, speech rate)
- **Qwen2.5-3B:** 4-bit quantized LLM for semantic evaluation

7 Results

7.1 Output Artifacts

- **timeline.json**: Master timebase synchronization
- **qa_pairs.json**: Question-answer mappings with timing
- **candidate_behavior_metrics.json**: Confidence, fluency metrics
- **relevance_scores.json**: LLM-evaluated relevance per Q&A
- **candidate_score_timeline.json**: Time-evolving performance scores

8 Evaluation and Validation

- Verdict categories: STRONG_HIRE, HIRE, BORDERLINE, NO_HIRE
- 5 competency dimensions: Technical depth, System design, Production experience, Communication clarity, Problem solving
- Matched keywords extraction from job description
- Incremental verdict progression throughout interview
- Confidence level provided with final recommendation

9 Future Enhancements

- Dashboard visualization for result presentation
- Real-time streaming analysis capability
- Multi-language support for transcription
- Enhanced video analytics (micro-expressions, body language)
- Integration with existing ATS platforms

10 Conclusion

- TIPS provides objective, data-driven candidate assessment
- Multi-stage pipeline enables comprehensive interview analysis
- LLM integration allows semantic evaluation against job requirements
- Temporal tracking shows candidate performance evolution
- Modular design facilitates future enhancements and optimizations